

Networking for DevOps

Fundamentals to Azure Cloud Networking

Complete Study Guide

OSI Model | TCP/UDP | IP & CIDR | DNS | HTTP | Ports | Azure Networking

DevOps Engineer Job Preparation

TABLE OF CONTENTS

- PART 1: NETWORKING FUNDAMENTALS
- PART 2: AZURE CLOUD NETWORKING
- PART 3: INTERVIEW PREPARATION
- PART 4: HANDS-ON PRACTICE & PROJECTS
- PART 5: COMMON MISTAKES & TROUBLESHOOTING
- PART 6: QUICK REVISION GUIDE

PART 1

NETWORKING FUNDAMENTALS

1. OSI Model (Open Systems Interconnection)

1.1 What is the OSI Model?

The OSI model is a conceptual framework that standardizes how different network protocols and systems communicate. It divides networking into 7 layers, each with specific responsibilities. Think of it as a blueprint for how data travels from your application to the network and back.

Why DevOps Engineers Need to Know This: When you're troubleshooting why a container can't talk to a database, or why your load balancer isn't working, understanding OSI layers helps you quickly identify where the problem is — application layer? Network layer? Transport layer?

The 7 Layers (Remember: Please Do Not Throw Sausage Pizza Away)

| Layer | Name | Protocols | Responsibility | Examples |
|-------|--------------|--------------------------------------|--|--|
| 7 | Application | HTTP, HTTPS, FTP, DNS, SMTP, SSH | What users interact with directly | Web browsers, email clients, REST APIs |
| 6 | Presentation | SSL/TLS, encryption, data formatting | Translates data formats, encryption | JPEG, PNG, encryption of HTTPS |
| 5 | Session | NetBIOS, RPC, session management | Establishes, maintains, terminates connections | Login sessions, database connections |
| 4 | Transport | TCP, UDP | Reliable/unreliable data transfer, ports | Port 443 for HTTPS, Port 22 for SSH |
| 3 | Network | IP, ICMP, routing | Logical addressing, routing between networks | Routers, IP addresses (192.168.1.1) |
| 2 | Data Link | Ethernet, Wi-Fi, MAC addresses | Physical addressing within local network | Switches, MAC addresses, VLANs |
| 1 | Physical | Cables, Wi-Fi signals, NICs | Actual physical connection | Ethernet cables, fiber optics, radio waves |

Real-World DevOps Example

When you type '<https://api.company.com/users>' in your browser and hit enter:

- Layer 7 (Application): Browser creates HTTP GET request
- Layer 6 (Presentation): SSL/TLS encrypts the request
- Layer 5 (Session): Establishes session with the server
- Layer 4 (Transport): TCP breaks data into segments, adds port 443
- Layer 3 (Network): IP adds source/destination IP addresses for routing
- Layer 2 (Data Link): Ethernet adds MAC addresses for local network
- Layer 1 (Physical): Data converted to electrical signals on the wire

The server receives it in reverse order (1→7) and sends response back the same way.

Troubleshooting by Layer (DevOps Application)

| Layer | Symptom | What to Check |
|---------|----------------------------|---|
| Layer 7 | Application not responding | Check: app logs, HTTP status codes, API timeouts |
| Layer 4 | Connection refused | Check: is port open? Firewall blocking? Service running? |
| Layer 3 | Cannot ping server | Check: IP address correct? Routing tables? VPN connected? |
| Layer 2 | No network access | Check: Ethernet cable plugged in? WiFi connected? MAC address filtered? |
| Layer 1 | Physical connection issue | Check: Cable damaged? Network card working? Port disabled? |

★ NOTE: For DevOps, focus on Layers 3, 4, and 7. These are where 90% of issues occur: IP/routing problems (L3), port/firewall issues (L4), and application errors (L7).

2. TCP vs UDP

2.1 Understanding Transport Layer Protocols

TCP and UDP are both Layer 4 (Transport) protocols. They both move data between applications on different machines, but they work very differently. Choosing the wrong one can break your application or waste resources.

TCP (Transmission Control Protocol)

TCP is connection-oriented and reliable. It guarantees that data arrives in order and without errors. Think of it like a phone call — you establish connection, talk, and hang up.

► How TCP Works

Three-Way Handshake (Connection Establishment):

Client sends SYN (synchronize) to server

Server responds with SYN-ACK (acknowledge)

Client sends ACK back

Connection established — now data can flow

Data Transfer with Acknowledgment:

Every packet sent must be acknowledged by receiver

If acknowledgment not received, packet is retransmitted

Packets are numbered, so receiver can reorder them if needed

Four-Way Handshake (Connection Termination):

Both sides send FIN (finish) and ACK to close gracefully

► TCP Features

Reliable — guaranteed delivery, packets retransmitted if lost

Ordered — data arrives in the same order it was sent

Error-checked — checksum detects corrupted packets

Flow control — sender doesn't overwhelm receiver

Congestion control — adjusts sending rate based on network conditions

UDP (User Datagram Protocol)

UDP is connectionless and unreliable (but fast!). It sends data without establishing a connection or confirming delivery. Think of it like sending postcards — you drop them in the mail and hope they arrive.

► How UDP Works

No connection setup — just send packets immediately

No acknowledgment — sender doesn't know if packet arrived

No guarantee of order — packets may arrive out of sequence

No retransmission — lost packets are lost forever

► UDP Features

Fast — no handshake overhead, no waiting for acknowledgments

Lightweight — smaller packet headers (8 bytes vs TCP's 20 bytes)

No connection state — server doesn't track connections, scales better

Broadcast/Multicast support — send one packet to multiple receivers

TCP vs UDP — When to Use Which

| Feature | TCP | UDP |
|----------------|---|--|
| Connection | Connection-oriented (3-way handshake) | Connectionless (no handshake) |
| Reliability | Guaranteed delivery, retransmits lost packets | No guarantee, packets can be lost |
| Ordering | Data arrives in order | Packets may arrive out of order |
| Speed | Slower (acknowledgment overhead) | Faster (no overhead) |
| Error Checking | Extensive | Basic checksum only |
| Use Cases | HTTP, HTTPS, FTP, SSH, databases, email | DNS, video streaming, VoIP, gaming, DHCP |
| When to Use | When you MUST have reliable delivery | When speed matters more than reliability |

Real DevOps Examples

► Example 1: Why DNS Uses UDP

DNS queries are tiny (usually <512 bytes) and need to be fast. If a DNS packet is lost, the client just resends the query — much simpler than maintaining TCP connections for millions of DNS queries per second. However, for large DNS responses (like zone transfers), TCP is used.

► Example 2: Why Databases Use TCP

When your application queries a database, you absolutely need every byte of data. If even one character is missing from a SQL query result, the application will break. TCP's guaranteed delivery is essential. PostgreSQL, MySQL, MongoDB all use TCP.

► Example 3: Why Live Video Streaming Uses UDP

In a live stream or video call, if a packet is lost, retransmitting it 3 seconds later is useless — the moment has passed. Better to skip that frame and continue. UDP's speed matters more than perfection. Netflix, Zoom, Twitch all use UDP-based protocols.

💡 TIP: In interviews, if asked 'Should I use TCP or UDP?', think: 'Can I tolerate data loss?' If no → TCP. If speed is critical and loss is acceptable → UDP.

3. IP Addressing & CIDR

3.1 IP Addresses

An IP address uniquely identifies a device on a network. IPv4 uses 32-bit addresses (4 numbers from 0-255). IPv6 uses 128-bit addresses (we'll focus on IPv4 as it's still dominant).

IPv4 Address Structure

192.168.1.100

This is four octets (8 bits each) separated by dots. Each octet ranges from 0 to 255.

Binary: 11000000.10101000.00000001.01100100

Decimal: 192 .168 .1 .100

Public vs Private IP Addresses

| Type | Visibility | Characteristics | Examples |
|------------|---------------------------|--|--|
| Public IP | Routable on the internet | Assigned by ISP, globally unique | 8.8.8.8, 40.112.72.205 |
| Private IP | Only within local network | NOT routable on internet, reused in different networks | 10.x.x.x, 172.16-31.x.x, 192.168.x.x |

Private IP Ranges (RFC 1918)

10.0.0.0 to 10.255.255.255 (10.0.0.0/8) — 16 million addresses

172.16.0.0 to 172.31.255.255 (172.16.0.0/12) — 1 million addresses

192.168.0.0 to 192.168.255.255 (192.168.0.0/16) — 65,536 addresses

★ NOTE: Your home router likely uses 192.168.1.x. Azure VNets commonly use 10.0.0.0/16 or 172.16.0.0/16.

Special IP Addresses

| Address | Purpose |
|-----------|---|
| 127.0.0.1 | localhost / Loopback — always refers to current machine |

| | |
|-----------------|--|
| 0.0.0.0 | All interfaces / default route in routing tables |
| 255.255.255.255 | Broadcast address — send to all devices on local network |
| 169.254.x.x | APIPA (Automatic Private IP) — self-assigned when DHCP fails |

3.2 CIDR (Classless Inter-Domain Routing)

CIDR notation (e.g., 192.168.1.0/24) defines an IP range using a network prefix. It replaced the old 'Class A/B/C' system and is essential for subnetting in cloud environments.

Understanding CIDR Notation

192.168.1.0/24

192.168.1.0 = network address (first address in range)

/24 = subnet mask — first 24 bits are network, last 8 bits are hosts

Subnet mask: 255.255.255.0 (binary: 11111111.11111111.11111111.00000000)

Number of addresses = $2^{(32-24)} = 2^8 = 256$ addresses

Usable addresses = 256 - 2 = 254 (first is network, last is broadcast)

Address range: 192.168.1.0 to 192.168.1.255

Common CIDR Blocks

| CIDR | Subnet Mask | Total IPs | Usable Hosts | Typical Use |
|------|-----------------|------------|--------------------|----------------------------------|
| /32 | 255.255.255.255 | 1 | 1 (single host) | Firewall rule for specific IP |
| /31 | 255.255.255.254 | 2 | 2 (point-to-point) | Router-to-router links |
| /30 | 255.255.255.252 | 4 | 2 | Small point-to-point connections |
| /29 | 255.255.255.248 | 8 | 6 | Very small subnet |
| /28 | 255.255.255.240 | 16 | 14 | Small office network |
| /27 | 255.255.255.224 | 32 | 30 | Small subnet |
| /26 | 255.255.255.192 | 64 | 62 | Medium subnet |
| /25 | 255.255.255.128 | 128 | 126 | Medium-large subnet |
| /24 | 255.255.255.0 | 256 | 254 | Standard subnet (254 hosts) |
| /23 | 255.255.254.0 | 512 | 510 | Larger subnet |
| /22 | 255.255.252.0 | 1,024 | 1,022 | Large subnet |
| /21 | 255.255.248.0 | 2,048 | 2,046 | Very large subnet |
| /20 | 255.255.240.0 | 4,096 | 4,094 | Huge subnet |
| /16 | 255.255.0.0 | 65,536 | 65,534 | Azure VNet default size |
| /8 | 255.0.0.0 | 16,777,216 | 16,777,214 | Entire Class A network |

CIDR Calculation Examples

► **Example 1: 10.0.0.0/16**

Network: 10.0.0.0

Subnet mask: 255.255.0.0

First 16 bits locked, last 16 bits available for hosts

Total addresses: $2^{16} = 65,536$

Range: 10.0.0.0 to 10.0.255.255

Usable: 10.0.0.1 to 10.0.255.254 (65,534 hosts)

Common for: Azure VNet address space

► **Example 2: 172.16.5.0/24**

Network: 172.16.5.0

Subnet mask: 255.255.255.0

Range: 172.16.5.0 to 172.16.5.255

Usable: 172.16.5.1 to 172.16.5.254 (254 hosts)

Common for: Azure subnet within VNet

Subnetting in Azure (Real Example)

You have Azure VNet with 10.0.0.0/16. You want to create multiple subnets:

| | | |
|-------|-------------|----------------------------|
| VNet: | 10.0.0.0/16 | (10.0.0.0 to 10.0.255.255) |
|-------|-------------|----------------------------|

| | | |
|------------------|-------------|---------------------------------------|
| Subnet 1 (Web) : | 10.0.1.0/24 | (10.0.1.0 to 10.0.1.255) – 254 IPs |
|------------------|-------------|---------------------------------------|

| | | |
|------------------|-------------|---------------------------------------|
| Subnet 2 (App) : | 10.0.2.0/24 | (10.0.2.0 to 10.0.2.255) – 254 IPs |
|------------------|-------------|---------------------------------------|

| | | |
|-----------------|-------------|---------------------------------------|
| Subnet 3 (DB) : | 10.0.3.0/24 | (10.0.3.0 to 10.0.3.255) – 254 IPs |
|-----------------|-------------|---------------------------------------|

| | | |
|----------------------|-------------|-------------------------------------|
| Subnet 4 (Gateway) : | 10.0.4.0/27 | (10.0.4.0 to 10.0.4.31) – 30 IPs |
|----------------------|-------------|-------------------------------------|

NOTE: Azure reserves 5 IP addresses in each subnet (first 4 and last 1). So a /24 subnet gives you $254 - 5 = 249$ usable IPs.

TIP: For interviews, memorize: /16 \approx 65k IPs, /24 = 256 IPs, /32 = single IP.

4. DNS (Domain Name System)

4.1 What is DNS?

DNS is the internet's phone book. It translates human-readable domain names (google.com) into machine-readable IP addresses (142.250.185.46). Without DNS, you'd have to memorize IP addresses for every website.

How DNS Works (Step-by-Step)

User types 'www.company.com' in browser

Browser checks local cache — is IP already known?

If not cached, query sent to recursive DNS resolver (usually ISP's DNS or 8.8.8.8)

Resolver checks its cache. If not found, queries root DNS servers

Root server responds: 'For .com domains, ask the .com TLD servers'

Resolver queries .com TLD server: 'Where is company.com?'

TLD responds: 'Ask the authoritative name server for company.com'

Resolver queries authoritative server for company.com

Authoritative server responds with IP: 40.112.72.205

Resolver caches the result and returns IP to browser

Browser connects to 40.112.72.205

All this happens in milliseconds!

DNS Record Types

| Record Type | Purpose | Example |
|------------------------|---------------------------------------|--|
| A (Address) | Maps domain to IPv4 address | www.company.com → 40.112.72.205 |
| AAAA (Quad-A) | Maps domain to IPv6 address | www.company.com → 2001:db8::1 |
| CNAME (Canonical Name) | Maps domain to another domain (alias) | blog.company.com → company.github.io |
| MX (Mail Exchange) | Specifies mail servers for domain | company.com → mail.google.com (priority 10) |
| TXT (Text) | Stores text data for various purposes | SPF records, domain verification, DKIM |
| NS (Name Server) | Specifies authoritative name servers | company.com → ns1.azure-dns.com |
| PTR (Pointer) | Reverse DNS — maps IP to domain name | 40.112.72.205 → server.company.com |

| | | |
|--------------------------|--------------------------------|--|
| SRV (Service) | Specifies location of services | _sip._tcp.company.com → sipserver.com:5060 |
| SOA (Start of Authority) | Administrative info about zone | Primary name server, admin email, TTL |

Real DevOps Examples

► Example 1: A Record for Web Server

```
# DNS Zone: company.com

www      A      40.112.72.205    # Web server IP
api      A      40.112.73.120    # API server IP
@       A      40.112.72.205    # Root domain points to web server
```

► Example 2: CNAME for Subdomains

```
# Point multiple subdomains to same Azure App Service

www      CNAME   myapp.azurewebsites.net
app      CNAME   myapp.azurewebsites.net
staging  CNAME   myapp-staging.azurewebsites.net
```

► Example 3: MX Records for Email

```
# Route email to Microsoft 365

@  MX  10  company-com.mail.protection.outlook.com
@  MX  20  backup-mx.company.com
```

Lower priority number = preferred mail server

► Example 4: TXT for Domain Verification

```
# Prove you own the domain to Azure

@  TXT  MS=ms12345678

# SPF record to prevent email spoofing
@  TXT  v=spf1 include:_spf.microsoft.com ~all
```

DNS TTL (Time To Live)

TTL tells DNS resolvers how long to cache the record (in seconds).

| TTL | Use Case | Example |
|---------------|-----------------------------|---------------------------------------|
| 3600 (1 hour) | Standard for stable records | Production servers that rarely change |

| | | |
|------------------|-------------------------------|--|
| 300 (5 minutes) | For records that might change | During migration or testing |
| 60 (1 minute) | For active DNS changes | Blue-green deployments, failover testing |
| 86400 (24 hours) | For very stable records | Name servers, rarely-changing infrastructure |

 **NOTE:** Before making DNS changes in production, lower TTL to 300 (5 min) a day in advance. After change, raise it back to 3600 to reduce query load.

Public DNS Resolvers (Memorize These)

| | |
|--------------------|---|
| Google Public DNS: | 8.8.8.8 and 8.8.4.4 |
| Cloudflare DNS: | 1.1.1.1 and 1.0.0.1 |
| Quad9: | 9.9.9.9 |
| Azure DNS: | 168.63.129.16 (internal Azure VNet DNS) |

 **TIP:** If DNS is not resolving, test with: nslookup google.com 8.8.8.8

5. HTTP Status Codes

5.1 Understanding HTTP Status Codes

HTTP status codes are 3-digit numbers returned by web servers to indicate the result of a request. As a DevOps engineer, you'll see these in logs constantly and need to troubleshoot based on them.

Status Code Categories

| Range | Category | Meaning | Who's Responsible |
|-------|---------------|---|-------------------------|
| 1xx | Informational | Request received, continuing process | Rarely seen in practice |
| 2xx | Success | Request successfully received, understood, and accepted | Everything worked |
| 3xx | Redirection | Further action needed to complete request | Resource moved |
| 4xx | Client Error | Client sent a bad request | User's fault |
| 5xx | Server Error | Server failed to fulfill valid request | Server's fault |

Most Important Status Codes (Memorize These)

► 2xx — Success

| Code | Meaning | Example |
|----------------|---|--|
| 200 OK | Request succeeded. Most common success code. | GET /api/users → returns user list |
| 201 Created | Resource was successfully created. | POST /api/users → new user created |
| 202 Accepted | Request accepted for processing but not complete yet. | Async job submitted |
| 204 No Content | Success but no content to return. | DELETE /api/users/123 → success, nothing to return |

► 3xx — Redirection

| Code | Meaning | Example |
|------------------------|--|---|
| 301 Moved Permanently | Resource permanently moved to new URL. Update bookmarks! | http://site.com → https://site.com |
| 302 Found | Temporary redirect. Original URL still valid. | Maintenance redirect |
| 304 Not Modified | Cached version is still valid. Don't download again. | Browser cache is fresh |
| 307 Temporary Redirect | Like 302 but preserves HTTP method (POST stays POST). | Temporary failover |

► 4xx — Client Errors

| Code | Meaning | Example |
|------------------------|--|----------------------------------|
| 400 Bad Request | Invalid request syntax, malformed JSON, missing fields. | POST with invalid JSON body |
| 401 Unauthorized | Authentication required. User not logged in. | Accessing /admin without token |
| 403 Forbidden | Authenticated but not authorized. You don't have permission. | Regular user accessing admin API |
| 404 Not Found | Resource doesn't exist. Wrong URL or resource deleted. | GET /api/users/999999 |
| 405 Method Not Allowed | HTTP method not supported on this endpoint. | POST to read-only endpoint |
| 408 Request Timeout | Server timed out waiting for request. | Slow client connection |
| 429 Too Many Requests | Rate limit exceeded. Too many requests too fast. | API throttling |

► 5xx — Server Errors

| Code | Meaning | Example |
|---------------------------|---|--|
| 500 Internal Server Error | Generic server error. Something crashed. | Unhandled exception in code |
| 502 Bad Gateway | Upstream server returned invalid response. Gateway/proxy can't reach app. | App server down, load balancer can't connect |
| 503 Service Unavailable | Server overloaded or down for maintenance. | Deployment in progress, server restarting |

| | | |
|---------------------|---|---|
| 504 Gateway Timeout | Upstream server didn't respond in time. Slow backend. | Database query took 60 seconds, timeout at 30 |
|---------------------|---|---|

Real DevOps Scenarios

► Scenario 1: 502 Bad Gateway During Deployment

You deploy a new version. Users immediately get 502 errors.

Cause: Application hasn't started yet, or crashed on startup. Load balancer forwards requests but app isn't listening.

Fix: Check app logs. Is app running? Port correct? Health check endpoint responding?

► Scenario 2: 504 Gateway Timeout on API

API works fine usually, but sometimes returns 504.

Cause: Database query is slow or times out. Application Gateway timeout (default 30 seconds) is reached.

Fix: Optimize slow queries. Add database indexes. Increase Gateway timeout if queries legitimately take longer.

► Scenario 3: 429 Too Many Requests

Mobile app users complain they can't log in during peak hours.

Cause: API rate limiting kicked in. Too many requests from same IP/user.

Fix: Increase rate limits if legitimate traffic. Add caching to reduce backend calls. Implement exponential backoff in client.

💡 TIP: In interviews, if asked 'What's the difference between 401 and 403?', answer: '401 means NOT logged in (no credentials or invalid credentials). 403 means logged in but FORBIDDEN (not authorized for this resource).'

6. Important Port Numbers

6.1 Well-Known Ports (0-1023)

Well-known ports are standardized and used by common services. You MUST memorize these for DevOps interviews and daily troubleshooting.

Essential Ports (Memorize These!)

| Port | Service | Description | Common Use |
|-------|------------|--|---|
| 20/21 | FTP | File Transfer Protocol (data/control) | Legacy file transfer |
| 22 | SSH | Secure Shell — remote server access | ssh user@server |
| 23 | Telnet | Unencrypted remote access (NEVER use!) | Replaced by SSH |
| 25 | SMTP | Email sending (Simple Mail Transfer Protocol) | Mail servers |
| 53 | DNS | Domain Name System (UDP primary, TCP for zone transfers) | DNS queries |
| 80 | HTTP | Unencrypted web traffic | http://website.com |
| 110 | POP3 | Email retrieval (Post Office Protocol) | Older email protocol |
| 143 | IMAP | Email retrieval (better than POP3) | Modern email |
| 443 | HTTPS | Encrypted web traffic (HTTP over TLS) | https://website.com |
| 445 | SMB | Windows file sharing (Server Message Block) | File shares |
| 3306 | MySQL | MySQL database default port | Database connections |
| 3389 | RDP | Remote Desktop Protocol (Windows) | Remote into Windows VM |
| 5432 | PostgreSQL | PostgreSQL database | Database connections |
| 6379 | Redis | Redis cache and message broker | Caching layer |
| 8080 | HTTP-Alt | Alternative HTTP (Tomcat, Jenkins, test servers) | Dev/test environments |
| 27017 | MongoDB | MongoDB NoSQL database | Database connections |

Kubernetes & Container Ports

| Port(s) | Service | Description | Usage |
|-------------|----------------|-----------------------------------|-----------------------------|
| 2379/2380 | etcd | Kubernetes cluster state database | K8s control plane |
| 6443 | Kubernetes API | API server secure port | kubectl commands |
| 10250 | kubelet | Kubelet API (on each node) | Node management |
| 30000-32767 | NodePort | K8s NodePort service range | External access to services |

Azure-Specific Ports

| Port | Service | Usage | Azure Context |
|-----------|------------|-------------------------------|----------------------|
| 3389 | RDP | Remote Desktop to Windows VMs | Windows VM access |
| 22 | SSH | SSH to Linux VMs | Linux VM access |
| 1433 | SQL Server | Azure SQL Database | Database connections |
| 5671/5672 | AMQP | Azure Service Bus, Event Hubs | Message queuing |

DevOps Tools Ports

| Port | Tool | Purpose | DevOps Use |
|-----------|---------------|--|----------------------|
| 8080 | Jenkins | Jenkins web UI and API | CI/CD server |
| 9090 | Prometheus | Prometheus metrics server | Monitoring |
| 3000 | Grafana | Grafana dashboards | Visualization |
| 9200/9300 | Elasticsearch | Elasticsearch HTTP API / cluster communication | Log aggregation |
| 5601 | Kibana | Kibana web UI | Log visualization |
| 2375/2376 | Docker | Docker daemon API (unsecured/TLS) | Container management |

Security Note on Ports

⚠️ WARNING: NEVER expose ports like 22 (SSH), 3389 (RDP), 3306 (MySQL) directly to the internet. Use VPN, bastion hosts, or Azure Bastion instead. Bots constantly scan for these open ports.

Real DevOps Example: Firewall Rule

```
# Azure NSG rule to allow HTTPS from internet
Priority: 100
Source: Internet (or specific IP)
Destination: Any
Port: 443
Protocol: TCP
Action: Allow
```

```
# Allow SSH only from corporate VPN
Priority: 200
Source: 203.0.113.0/24 (corporate network)
Destination: Any
Port: 22
Protocol: TCP
Action: Allow
```

💡 TIP: To test if a port is open: telnet server-ip 443 or nc -zv server-ip 443

PART 2

AZURE CLOUD NETWORKING

7. Azure Virtual Network (VNet)

7.1 What is Azure VNet?

Azure Virtual Network (VNet) is your private network in the cloud. It's logically isolated from other VNets and the internet. Think of it as your own data center network in Azure — you define the IP address space, create subnets, and control traffic flow.

Key Concepts

Isolation — VNet is private and isolated by default. Resources in different VNets cannot communicate unless explicitly connected.

Address Space — Define your own IP ranges (e.g., 10.0.0.0/16) using private IPs from RFC 1918.

Subnets — Divide VNet into smaller segments for organization and security.

Region-Specific — VNet exists in one Azure region, but can connect to VNets in other regions via peering.

Creating a VNet (Conceptual Steps)

Choose address space (e.g., 10.0.0.0/16 gives 65,536 IPs)

Create subnets within that address space:

Frontend subnet: 10.0.1.0/24 (254 usable IPs)

Backend subnet: 10.0.2.0/24

Database subnet: 10.0.3.0/24

Deploy resources (VMs, App Services, AKS) into appropriate subnets

Configure Network Security Groups (NSGs) for firewall rules

Real DevOps Example: Multi-Tier Application VNet

```
VNet Name: ProductionVNet  
Address Space: 10.1.0.0/16
```

Subnets:

- WebTier: 10.1.1.0/24 (Web servers, load balancer)
- AppTier: 10.1.2.0/24 (Application servers, APIs)
- DataTier: 10.1.3.0/24 (Databases, cache)
- GatewaySubnet: 10.1.4.0/27 (VPN Gateway – special subnet name required)
- AzureBastionSubnet: 10.1.5.0/27 (Azure Bastion – special name required)

⚠ NOTE: GatewaySubnet and AzureBastionSubnet are special subnet names required by Azure. Don't rename them.

7.2 Subnets

Subnets divide a VNet into smaller networks. Each subnet is a range of IP addresses within the VNet address space.

Why Use Subnets?

Security Segmentation — Separate web, app, and database tiers with different security rules

Organization — Group resources by function, environment, or team

NSG Application — Apply different firewall rules to different subnets

Routing Control — Route traffic differently based on subnet

Azure Reserved IPs in Subnets

Azure reserves 5 IP addresses in EVERY subnet:

| IP Address | Purpose |
|------------|---------------------------------------|
| x.x.x.0 | Network address (not usable) |
| x.x.x.1 | Reserved by Azure for default gateway |
| x.x.x.2 | Reserved by Azure for DNS |
| x.x.x.3 | Reserved by Azure for DNS |
| x.x.x.255 | Broadcast address (not usable) |

So a /24 subnet (256 IPs) has only 251 usable IPs, not 256.

Subnet Sizing Guidelines

| CIDR | Total IPs | Usable (minus Azure reserved 5) | Typical Use |
|----------------|-------------|---|-------------|
| /28 (16 IPs) | 11 usable | Small subnet for gateways, bastion | |
| /27 (32 IPs) | 27 usable | GatewaySubnet, BastionSubnet minimum size | |
| /26 (64 IPs) | 59 usable | Small application tier | |
| /25 (128 IPs) | 123 usable | Medium subnet | |
| /24 (256 IPs) | 251 usable | Standard subnet — good for most cases | |
| /23 (512 IPs) | 507 usable | Large subnet for scale-out apps | |
| /22 (1024 IPs) | 1019 usable | Very large subnet or AKS clusters | |

TIP: For AKS clusters, plan for growth. A /24 subnet supports ~30 nodes. Use /22 or /20 for production AKS.

8. Network Security Groups (NSG)

8.1 What is an NSG?

Network Security Group is Azure's distributed firewall. It filters network traffic to and from Azure resources based on rules you define. Think of it as iptables or firewalld in Linux, but managed by Azure.

How NSG Works

Rule-Based Filtering — Define allow/deny rules based on source, destination, port, protocol

Priority-Based Evaluation — Rules processed in priority order (100-4096, lower number = higher priority)

Stateful — If you allow inbound on port 443, return traffic is automatically allowed

Applied to Subnet or NIC — Can attach NSG to subnet (affects all resources) or individual network interface (affects one VM)

NSG Rule Structure

| Component | Description |
|-------------|---|
| Priority | 100-4096 (lower = higher priority). Process in order until match found. |
| Name | Descriptive name for the rule (e.g., 'AllowHTTPS') |
| Port | Single port (443), range (8000-9000), or * (all) |
| Protocol | TCP, UDP, ICMP, Any |
| Source | IP address, CIDR range, Service Tag (Internet, VirtualNetwork, AzureLoadBalancer) |
| Destination | IP address, CIDR range, Service Tag |
| Action | Allow or Deny |
| Direction | Inbound (traffic coming TO resource) or Outbound (traffic leaving FROM resource) |

Example NSG Rules

► Web Server NSG (Applied to WebTier Subnet)

| Priority | Name | Port | Protocol | Source | Destination | Action | Direction |
|----------|-----------------|------|----------|----------------|-------------|---------|-----------|
| 100 | AllowHTTPS | 443 | TCP | Internet | Any | Allow | Inbound |
| 110 | AllowHTTP | 80 | TCP | Internet | Any | Allow | Inbound |
| 200 | AllowSSHFromVPN | 22 | TCP | 203.0.113.0/24 | Any | Allow | Inbound |
| 300 | DenyAllInbound | * | * | * | Deny | Inbound | * |

► Database Subnet NSG

| Priority | Name | Port | Protocol | Source | Destination | Action | Direction |
|----------|------------------|------|----------|-------------|-------------|--------|-----------|
| 100 | AllowFromAppTier | 5432 | TCP | 10.1.2.0/24 | Any | Allow | Inbound |
| 200 | DenyAllElse | * | * | * | * | Deny | Inbound |

This ensures database only accepts connections from the App subnet, nowhere else.

Default NSG Rules (Cannot Delete)

Every NSG has default rules at the bottom (priority 65000+). You can't delete them, but your custom rules override them.

| Default Rule | Priority | Action | Meaning |
|-------------------------------|----------|--------|--|
| AllowVNetInbound | 65000 | Allow | Traffic within VNet allowed by default |
| AllowAzureLoadBalancerInbound | 65001 | Allow | Azure Load Balancer health probes |

| | | | |
|----------------|-------|------|---|
| DenyAllInbound | 65500 | Deny | Deny everything else (unless your rules allow it) |
|----------------|-------|------|---|

Service Tags (Simplified Source/Destination)

Instead of specifying IP ranges, use service tags:

| Service Tag | What It Represents |
|-------------------|---|
| Internet | Any public internet IP address |
| VirtualNetwork | All IPs within the VNet and connected VNets |
| AzureLoadBalancer | Azure's load balancer infrastructure |
| Storage | Azure Storage service IP ranges |
| Sql | Azure SQL Database IP ranges |
| AzureCloud | All Azure datacenter IPs |

NSG Best Practices

Least Privilege — Start with deny-all, then allow only what's needed

Use Service Tags — Easier than maintaining IP lists

Document Rules — Use descriptive names like 'AllowHTTPSPFromCDN' not 'Rule1'

Subnet-Level NSGs — Apply to subnet for all resources, override with NIC-level NSG if needed

Audit Regularly — Review NSG rules quarterly, remove unused rules

⚠️ WARNING: Common mistake: Opening port 22 or 3389 to 0.0.0.0/0 (entire internet). This is scanned and attacked within minutes. Use Azure Bastion or VPN instead.

9. Load Balancer vs Application Gateway

9.1 When to Use Which?

Azure has TWO main load balancing services. Choosing the wrong one is a common mistake. Here's how to decide:

| Feature | Azure Load Balancer | Application Gateway |
|-----------------|------------------------------|--------------------------------------|
| OSI Layer | Layer 4 (Transport) | Layer 7 (Application) |
| Traffic Type | Any TCP/UDP traffic | HTTP/HTTPS only |
| Routing | Based on IP address and port | Based on URL path, hostname, headers |
| SSL Termination | No (passes through) | Yes (offloads SSL from backend) |

| | | |
|--------------------------|---|---|
| Web Application Firewall | No | Yes (optional WAF add-on) |
| WebSocket Support | Yes (generic TCP) | Yes (HTTP/1.1 upgrade) |
| Typical Use Case | Non-HTTP apps, VM load balancing | Web applications, API gateways, microservices |
| Cost | Lower | Higher |
| Example | Load balance database connections, RDP, SSH | Route /api to API servers, /images to static server |

9.2 Azure Load Balancer (Layer 4)

How It Works

Load Balancer distributes traffic based on IP address and port. It doesn't understand HTTP — it just forwards TCP/UDP packets.

- Client connects to Load Balancer public IP
- Load Balancer picks a backend VM based on algorithm (hash, round-robin)
- Traffic forwarded to VM private IP
- VM responds directly to client (or through LB depending on type)

Load Balancer Types

Public Load Balancer — Has public IP. Routes internet traffic to VMs in backend pool.

Internal Load Balancer — Has private IP. Routes traffic within VNet (e.g., app tier → database cluster).

Load Balancer Components

| Component | Purpose |
|---------------------|---|
| Frontend IP | The IP address clients connect to (public or private) |
| Backend Pool | VMs or VMSS instances that receive traffic |
| Health Probe | Checks if backend is healthy (TCP, HTTP). Unhealthy backends removed from pool. |
| Load Balancing Rule | Port mapping: frontend port → backend port, distribution algorithm |
| Inbound NAT Rule | Forward specific port to specific VM (e.g., SSH to individual VMs) |

Real Example: Database Load Balancer

Scenario: 3-node PostgreSQL cluster with read replicas

Internal Load Balancer:

Frontend IP: 10.1.3.100 (private IP in DB subnet)

```

Backend Pool: db-vm-1, db-vm-2, db-vm-3
Health Probe: TCP port 5432 every 5 seconds
LB Rule: 10.1.3.100:5432 → Backend:5432 (distribute connections)

```

```

Application servers connect to 10.1.3.100:5432
Load Balancer distributes to healthy database nodes

```

9.3 Application Gateway (Layer 7)

How It Works

Application Gateway understands HTTP. It can route based on URL paths, hostnames, headers. It terminates SSL and can inspect/modify HTTP requests.

Key Features

- URL-Based Routing** — Send /api/* to API servers, /images/* to static file servers
- Host-Based Routing** — Route app.company.com and api.company.com to different backend pools
- SSL Termination** — Handle SSL encryption/decryption at gateway, send plain HTTP to backends
- Web Application Firewall (WAF)** — Protect against OWASP Top 10 (SQL injection, XSS, etc.)
- Session Affinity** — Route same user to same backend (sticky sessions)
- Autoscaling** — Scale out/in based on traffic
- Redirection** — HTTP to HTTPS redirect built-in

Application Gateway Components

| Component | Purpose |
|-----------------|---|
| Frontend IP | Public and/or private IP where clients connect |
| Listener | Listens on specific port (80, 443) and hostname |
| Backend Pool | VMs, VMSS, App Services, or IP addresses |
| HTTP Settings | Port, protocol, cookie-based affinity, timeout |
| Routing Rule | Links listener → backend pool with path-based or host-based routing |
| Health Probe | Custom HTTP/HTTPS probe to check backend health |
| SSL Certificate | Upload cert for HTTPS termination |
| WAF Policy | Optional firewall rules to block malicious requests |

Real Example: Microservices Application Gateway

Frontend: appgateway.company.com (Public IP)

Routing Rules:

| | |
|-----------|---|
| / (root) | → Frontend Pool (React app VMs) |
| /api/* | → API Pool (Node.js VMs) |
| /auth/* | → Auth Pool (Authentication service) |
| /static/* | → Storage Account (Azure Blob with CDN) |

SSL Certificate: *.company.com

WAF: Enabled with OWASP 3.2 ruleset

Health Probe: GET /health every 30 seconds

When to Use Application Gateway

Web Applications — Multiple websites or services on same gateway

Microservices — Route different paths to different backend services

SSL Offload — Want to handle SSL at gateway, not at each backend

WAF Protection — Need to block SQL injection, XSS attacks

Cookie-Based Affinity — Stateful app needs same user → same server

When to Use Load Balancer

Non-HTTP Traffic — RDP, SSH, database connections, custom protocols

Simplicity & Cost — Basic TCP/UDP load balancing is all you need

Low Latency — Layer 4 has less overhead than Layer 7

Large Scale — Millions of connections, Load Balancer handles higher throughput

NOTE: You can use BOTH together! Application Gateway for web traffic → Load Balancer for backend database cluster.

10. VNet Peering

10.1 What is VNet Peering?

VNet Peering connects two Azure VNets so they can communicate privately as if they're on the same network. Traffic stays on Microsoft's backbone network — never goes over public internet.

Types of VNet Peering

VNet Peering (Regional) — Connect VNets in the same Azure region

Global VNet Peering — Connect VNets across different Azure regions

How VNet Peering Works

Create peering connection from VNet A to VNet B

Create reverse peering from VNet B to VNet A (peering is not automatic bi-directional!)

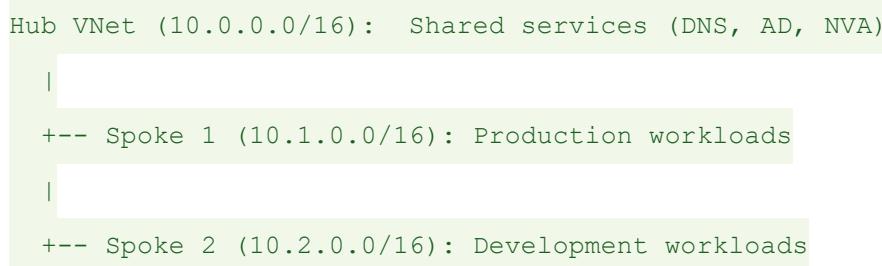
Resources in VNet A can now reach resources in VNet B using private IPs

No gateway needed, no encryption (but traffic never leaves Microsoft network)

VNet Peering Properties

| Feature | Details |
|-----------------------|--|
| Transitive | No. A↔B and B↔C does NOT mean A↔C. Must explicitly peer A↔C. |
| Bandwidth | No bandwidth limit — uses Azure backbone |
| Latency | Very low — direct connection |
| Address Space Overlap | NOT allowed. VNets must have non-overlapping IP ranges. |
| Cross-Region | Yes (Global VNet Peering) |
| Cross-Subscription | Yes (with proper permissions) |
| Cost | Ingress free, egress charged per GB (small cost) |

Real Example: Hub-and-Spoke Topology



```
+-- Spoke 3 (10.3.0.0/16) : Test workloads
```

Peering Configuration:

- Hub ↔ Spoke 1 (bidirectional peering)
- Hub ↔ Spoke 2 (bidirectional peering)
- Hub ↔ Spoke 3 (bidirectional peering)

Important: Spoke 1 cannot directly talk to Spoke 2. Traffic must go through Hub (use Network Virtual Appliance in Hub to route between spokes).

Common Use Cases

- Multi-Region Apps** — VNet in East US ↔ VNet in West Europe for DR/HA
- Hub-Spoke Architecture** — Centralized services in hub, isolated workloads in spokes
- Cross-Team Collaboration** — Development team VNet ↔ QA team VNet
- Hybrid Cloud** — On-prem connects to hub VNet, which peers to all spoke VNets

⚠️ WARNING: VNet address spaces cannot overlap! If VNet A is 10.0.0.0/16, VNet B cannot use 10.0.0.0/16 or any overlapping range.

11. Private Endpoints

11.1 What is a Private Endpoint?

Private Endpoint brings Azure PaaS services (Storage, SQL, Key Vault, etc.) into your VNet with a private IP address. Instead of accessing these services over the public internet, you access them via private IP within your VNet.

The Problem It Solves

Normally, Azure Storage is accessed via public endpoint: myaccount.blob.core.windows.net (public IP). This means:

Traffic goes over internet (even if within Azure)

Public endpoint is exposed (even with firewall rules)

Difficult to lock down for highly secure environments

Private Endpoint solves this:

Storage gets a private IP (e.g., 10.1.3.50) in your VNet

Access via private IP, traffic stays within VNet

Can disable public access entirely

How Private Endpoint Works

Create Private Endpoint in your VNet

Choose Azure service (Storage Account, SQL Database, Key Vault, etc.)

Private IP assigned from your VNet subnet

Private DNS zone created automatically (or you manage manually)

DNS resolution: myaccount.blob.core.windows.net → 10.1.3.50 (private IP)

Applications use same DNS name but connect via private IP

Supported Azure Services (Partial List)

Azure Storage (Blob, File, Queue, Table)

Azure SQL Database, Azure Database for PostgreSQL/MySQL

Azure Cosmos DB

Azure Key Vault

Azure Container Registry

Azure App Service, Azure Functions

Azure Event Hubs, Service Bus

Azure Cognitive Services

Real Example: Secure Database Access

Before Private Endpoint:

App in VNet connects to: mydb.database.windows.net

Resolves to: 40.112.72.205 (public IP)

Traffic: VNet → Internet → Azure SQL (even though both in Azure!)

Security: SQL firewall must allow VNet or specific IPs

After Private Endpoint:

Private Endpoint created in VNet subnet 10.1.3.0/24

Private IP assigned: 10.1.3.75

Private DNS: mydb.database.windows.net → 10.1.3.75

App connects to mydb.database.windows.net

Resolves to 10.1.3.75 (private IP)

Traffic: App → Private IP (stays in VNet)

Security: Public access disabled, only accessible from VNet

Private Endpoint vs Service Endpoint

| Feature | Private Endpoint | Service Endpoint |
|--------------|---|--|
| Traffic Path | Traffic stays within VNet | Traffic goes to public endpoint but via Azure backbone |
| IP Address | Service gets private IP in your VNet | Service keeps public IP |
| DNS | Private DNS zone resolves to private IP | Resolves to public IP (with source IP filtering) |
| Security | Can disable public access entirely | Public endpoint still exists (firewall rules required) |
| Cost | Charged per endpoint + data processed | Free (but limited to certain services) |
| Use Case | Maximum security, regulatory requirements | Simple VNet access, cost-sensitive |

💡 TIP: For production security: use Private Endpoint. For dev/test or cost savings: Service Endpoint is acceptable.

12. Azure Bastion

12.1 What is Azure Bastion?

Azure Bastion is a fully managed service that provides secure RDP/SSH access to your VMs without exposing them to the public internet. No need for public IPs on VMs, no need for VPN. Access VMs directly from Azure Portal over SSL.

The Problem It Solves

Traditional VM access methods and their problems:

| Method | Issues | Recommendation |
|-------------------------|---|-------------------------------|
| Public IP on VM + NSG | VM exposed to internet. Constant brute-force SSH/RDP attacks. Security risk. | ✗ Don't do this in production |
| Jump Box / Bastion Host | Dedicated VM with public IP for access. Still exposed, must maintain and patch. | ⚠ Better but still risky |
| VPN Gateway | Secure but requires client configuration, VPN software, certificates. | ✓ Good but complex |
| Azure Bastion | Fully managed, no public IPs on VMs, access via browser, no client software. | ✓ Best for most scenarios |

How Azure Bastion Works

Deploy Azure Bastion in your VNet (requires dedicated subnet named 'AzureBastionSubnet')

Bastion gets a public IP (but your VMs don't!)

In Azure Portal, click 'Connect' on a VM → Choose 'Bastion'

Enter VM credentials (username/password or SSH key)

Browser opens HTML5-based RDP or SSH session

Connection: Browser → Azure Bastion (HTTPS 443) → VM (RDP 3389 or SSH 22) over private IP

Azure Bastion Features

No Public IP on VMs — VMs remain fully private

No Client Software — Access from any browser, any device

Seamless Portal Integration — Click 'Connect' in Azure Portal, start session

SSL Protection — All traffic encrypted over HTTPS

Automatic Patching — Microsoft manages Bastion updates

No NSG Rules Needed on Bastion Subnet — Bastion handles required connectivity

Azure Bastion Requirements

Dedicated Subnet — Must be named 'AzureBastionSubnet', minimum /27 (32 IPs)

Standard Public IP — Bastion needs a Standard SKU public IP (but VMs don't!)

VNet Integration — Must be in same VNet as VMs you want to access

Real Example: Secure VM Access Architecture

VNet: ProductionVNet (10.1.0.0/16)

Subnets:

- AzureBastionSubnet: 10.1.0.0/27 (for Bastion service)
- AppServers: 10.1.1.0/24 (VMs with NO public IPs)
- DatabaseServers: 10.1.2.0/24 (VMs with NO public IPs)

Azure Bastion:

Name: ProductionBastion

Subnet: AzureBastionSubnet

Public IP: bastion-pip (only Bastion has this, VMs don't)

Access Flow:

Admin → Azure Portal → Click VM 'Connect' → Select Bastion

→ Browser opens SSH/RDP session

→ Traffic: Browser → Bastion (443) → VM (22/3389 private IP)

Bastion SKUs

| SKU | Scale | Features | Use Case |
|----------|------------------------------|--|----------------------|
| Basic | Up to 20 concurrent sessions | RDP/SSH access | Good for small teams |
| Standard | Up to 50 concurrent sessions | RDP/SSH + file upload/download, shareable link | Production workloads |

Cost Considerations

Azure Bastion is charged hourly plus data transfer. It's more expensive than a self-managed jump box VM, but far more secure and zero-maintenance.

| Option | Approx. Cost | Trade-off |
|---------------|--------------------|---------------------------------------|
| Azure Bastion | \$140/month + data | Fully managed, secure, no maintenance |
| Jump Box VM | \$30-70/month | Must maintain, patch, secure, monitor |

TIP: For production environments with compliance requirements (HIPAA, PCI-DSS), Azure Bastion is often mandatory — VMs cannot have public IPs.

 **WARNING:** Common mistake: Creating AzureBastionSubnet that's too small. Use /27 or /26 minimum. /28 will fail deployment.

PART 3

INTERVIEW QUESTIONS & ANSWERS

13. Interview Questions with Detailed Answers

13.1 Networking Fundamentals Questions

[BASIC] Explain the OSI model and which layers are most relevant for DevOps.

Answer: OSI model has 7 layers. For DevOps, the most important are: Layer 7 (Application) — where HTTP/HTTPS issues occur, troubleshooting API errors, status codes. Layer 4 (Transport) — TCP/UDP, port numbers, firewall rules, load balancing. Layer 3 (Network) — IP addressing, routing, VPN issues, VNet peering. When troubleshooting connectivity, I work from Layer 7 down: Is the app responding? Is the port open? Can I ping the IP? Is the cable plugged in?

[BASIC] What is the difference between TCP and UDP? Give real examples of when to use each.

Answer: TCP is connection-oriented and reliable — guarantees delivery and order. Use for: HTTP, databases, SSH, email — anything where data loss is unacceptable. UDP is connectionless and fast — no guarantee of delivery. Use for: DNS queries, video streaming, VoIP, gaming — where speed matters more than perfect reliability. Example: Database connections use TCP because missing one byte corrupts the entire query. Live video uses UDP because retransmitting a lost frame 3 seconds later is useless.

[INTERMEDIATE] Explain CIDR notation. What does 10.0.0.0/16 mean? How many usable IPs does it provide?

Answer: CIDR (Classless Inter-Domain Routing) defines a network range. 10.0.0.0/16 means: First 16 bits are network (10.0), last 16 bits are for hosts. Total IPs = $2^{(32-16)} = 65,536$. Usable IPs = 65,534 (first is network address, last is broadcast). In Azure, also subtract 5 reserved IPs per subnet. This notation is used everywhere in cloud: VNet address spaces, subnet ranges, firewall rules, routing tables.

[INTERMEDIATE] You see HTTP 502 Bad Gateway error. What does it mean and how do you troubleshoot?

Answer: 502 means the gateway/proxy received invalid response from upstream server. Common causes: (1) Application server is down or crashed — check if process is running. (2) Application is starting up but not ready yet — wait and retry. (3) Application responding on wrong port — verify app listens on port gateway expects. (4) Health check failing — check health endpoint returns 200. Troubleshooting steps: Check app logs for errors. Verify app process is running (systemctl status). Test app directly on backend (curl localhost:8080/health). Check gateway/load balancer health probe status.

[ADVANCED] Explain the difference between 401 Unauthorized and 403 Forbidden. Why are they often confused?

Answer: 401 Unauthorized means authentication failed or missing — you haven't proven who you are. User is anonymous or credentials are invalid. Client should send credentials (login). 403 Forbidden means you're authenticated (logged in) but not authorized for this resource — you don't have permission. Even with correct credentials, you're not allowed. They're confused because both prevent access, but 401 is 'who are you?' and 403 is 'I know who you are, but you can't do that'. Example: Accessing /admin without login → 401. Accessing /admin as regular user (not admin) → 403.

[ADVANCED] Design a subnet architecture for a 3-tier web application in Azure. Explain your choices.

Answer: I'd use a /16 VNet with multiple /24 subnets: VNet: 10.0.0.0/16. Subnets: (1) AzureBastionSubnet: 10.0.0.0/27 — for secure VM access, requires /27 minimum. (2) AppGatewaySubnet: 10.0.1.0/24 — Application Gateway, needs /24 for autoscaling. (3) WebTier: 10.0.10.0/24 — frontend VMs behind App Gateway. (4) AppTier: 10.0.20.0/24 — API/business logic servers. (5) DataTier: 10.0.30.0/24 — databases with strictest NSG rules. (6) GatewaySubnet: 10.0.100.0/27 — VPN Gateway for on-prem connectivity. NSG rules: Web tier allows 80/443 from internet via App Gateway. App tier allows app port only from Web tier. Data tier allows DB port only from App tier. Bastion subnet only allows 443 inbound.

13.2 Azure Networking Questions

[BASIC] What is Azure VNet and why do we need it?

Answer: Azure Virtual Network (VNet) is your private network in the cloud. It's logically isolated from other VNets and the internet. You need it to: (1) Deploy VMs, AKS, and other services securely. (2) Control network traffic with subnets and NSGs. (3) Connect to on-premises network via VPN or ExpressRoute. (4) Isolate workloads by environment (prod/dev/test) or team. Without VNet, resources would be on public internet directly, which is insecure. VNet gives you control over IP addressing, DNS,

routing, and security.

[BASIC] What is an NSG and how does it work?

Answer: Network Security Group (NSG) is Azure's firewall. It controls inbound and outbound traffic based on rules. Each rule has: priority (100-4096, lower = higher priority), source/destination (IP or service tag), port, protocol, and allow/deny action. NSG evaluates rules in priority order — first match wins. Can attach NSG to subnet (affects all resources in subnet) or network interface (affects individual VM). It's stateful — if you allow inbound 443, return traffic is automatic. Best practice: apply NSG to subnet for broad control, use NIC-level NSG for exceptions.

[INTERMEDIATE] What's the difference between Azure Load Balancer and Application Gateway?

Answer: Azure Load Balancer is Layer 4 (TCP/UDP) — routes based on IP and port. Doesn't understand HTTP. Fast, simple, low cost. Use for: database load balancing, non-HTTP apps, VM load balancing. Application Gateway is Layer 7 (HTTP/HTTPS only) — routes based on URL path, hostname, headers. Terminates SSL, includes WAF, supports cookie-based affinity. Use for: web apps, microservices, API gateways, anything HTTP. Example: Load Balancer for PostgreSQL cluster (TCP 5432). Application Gateway for routing /api to API servers and /app to frontend.

[INTERMEDIATE] Explain VNet Peering. Is it transitive?

Answer: VNet Peering connects two VNets privately over Microsoft's backbone. Resources communicate using private IPs. Benefits: no gateway needed, low latency, works across regions. Important: Peering is NOT transitive. If VNet A peers with B, and B peers with C, A cannot reach C directly — must explicitly peer A↔C. Address spaces cannot overlap — can't peer 10.0.0.0/16 with 10.0.0.0/16. Use cases: hub-and-spoke topology (hub has shared services, spokes are isolated workloads), multi-region DR, cross-team VNet access.

[ADVANCED] What is a Private Endpoint and when would you use it instead of Service Endpoint?

Answer: Private Endpoint brings Azure PaaS services (Storage, SQL, Key Vault) into your VNet with a private IP. Traffic never leaves VNet. Can disable public access entirely. Use when: maximum security required (compliance like HIPAA, PCI-DSS), want to block all public internet access, regulatory requirement for private connectivity. Service Endpoint is simpler/cheaper — traffic goes to public endpoint but via Azure backbone with source IP filtering. Still has public endpoint. Use when: basic VNet access is enough, cost is concern, security requirements are moderate. Key difference: Private Endpoint = service IN your VNet. Service Endpoint = VNet traffic TO public endpoint.

[ADVANCED] Design a secure network for VMs that require RDP/SSH access without exposing them to internet.

Answer: Best practice: Azure Bastion. Architecture: (1) Create dedicated

AzureBastionSubnet (/27 minimum) in VNet. (2) Deploy Azure Bastion service (gets public IP, but VMs don't). (3) VMs in private subnets with NO public IPs. (4) NSG on VM subnets allows RDP/SSH only from AzureBastionSubnet range. (5) Access: Portal → VM → Connect → Bastion → browser-based RDP/SSH over HTTPS. Benefits: VMs never exposed to internet (no brute-force attacks), no VPN client needed, access from any browser, fully managed (no patching), audit logs in Azure. Alternative for dev/test: self-managed jump box, but Bastion is production-grade.

13.3 Scenario-Based Questions

SCENARIO: Application Can't Connect to Database

Your web app in Azure suddenly can't connect to Azure SQL Database. It worked yesterday. Database is up and responding. How do you troubleshoot?

Expected Approach: Systematic approach: (1) Check app error logs — what's the exact error? Connection timeout? Auth failure? (2) Verify connection string — any recent changes? Is password correct? (3) Check Azure SQL firewall rules — is app's IP/VNet allowed? Check if 'Allow Azure Services' is enabled. (4) If using Private Endpoint — verify Private DNS resolution. Run nslookup from app server — does SQL FQDN resolve to private IP? (5) Check NSG rules — if database in VM, is port 1433 allowed from app subnet? (6) Test connectivity directly — from app VM: 'telnet sqlserver.database.windows.net 1433' or 'psping'. (7) Check SQL logs for blocked connections. (8) Verify service health in Azure Portal — any outages? Most common causes: NSG rule changed, SQL firewall rule removed, Private Endpoint DNS issue, expired password.

SCENARIO: High Latency Between Azure Regions

Your app in East US communicates with database in West Europe. Users complain of slow response times. How do you improve latency?

Expected Approach: Options to reduce cross-region latency: (1) Use Global VNet Peering instead of public internet — traffic on Microsoft backbone, much lower latency. (2) Deploy read replicas closer to users — Azure SQL active geo-replication, CosmosDB multi-region. Users read from nearest region. (3) Implement caching — Azure Cache for Redis in each region to avoid cross-region database calls. (4) Use CDN for static content — Azure Front Door or CDN caches content at edge, reduces backend calls. (5) Consider Azure Traffic Manager or Front Door for intelligent routing — route users to nearest healthy endpoint. (6) Evaluate if cross-region is necessary — can you replicate data asynchronously instead of synchronous calls? (7) If using ExpressRoute, verify BGP routing is optimal. (8) Monitor with Network Watcher — identify actual bottleneck (network vs application). Measure current latency before and after changes — use Connection Monitor.

SCENARIO: Intermittent Connection Timeouts

Users report your website times out intermittently — not always, maybe 1 in 20 requests. How do you diagnose this?

Expected Approach: Intermittent issues are hardest to debug. Approach: (1) Check Application Gateway / Load Balancer metrics — are backends failing health checks? Intermittent failures = likely one backend is unhealthy. (2) Review health probe logs — which backend is timing out? (3) Check that backend's logs during timeout windows — memory spikes? CPU at 100%? (4) Monitor NSG flow logs — are connections being blocked or dropped? (5) Check connection limits — is a backend hitting max connections? (6) Review autoscaling settings — are backends scaling fast enough during traffic spikes? (7) Check for network-level issues — Connection Monitor can detect packet loss. (8) If using App Service — check HTTP queue length metric, may need to scale out. (9) Database timeout? Check SQL connection pool exhaustion. (10) Enable verbose logging temporarily to catch next occurrence. Key: correlate timestamp of user complaint with backend metrics/logs to identify failing component.

SCENARIO: Cost Optimization for Networking

Your Azure bill shows high networking costs. What are common causes and how do you reduce them?

Expected Approach: Common high-cost scenarios: (1) VNet peering data transfer — especially Global VNet Peering across regions. Consider consolidating resources in fewer regions or using Private Endpoint instead of peering. (2) Public IP addresses — charged for Standard Public IPs. Audit unused IPs and delete. Use Azure Bastion instead of public IPs on each VM. (3) Application Gateway idle time — if low traffic, consider switching to Load Balancer (cheaper) or use App Gateway v2 autoscaling to scale to zero. (4) VPN Gateway always-on — if only needed occasionally, use Point-to-Site instead of Site-to-Site. (5) ExpressRoute unused bandwidth — right-size circuit, don't overprovision. (6) Outbound data transfer — largest cost driver. Reduce by: using CDN, caching aggressively, enabling compression, keeping data in same region. (7) NAT Gateway data processing — if very high traffic, verify it's needed. (8) Use Azure Cost Management to identify top networking resources by cost. Typical findings: 40% of networking cost is avoidable (unused IPs, over-provisioned gateways).

PART 4

HANDS-ON PRACTICE & PROJECTS

14. Practice Assignments & Projects

14.1 Networking Fundamentals — Practice

EASY Level

Calculate usable IPs for these CIDR blocks: /24, /26, /28, /30

Identify which of these are public vs private IPs: 192.168.1.1, 8.8.8.8, 10.0.0.1, 172.217.14.206

Look up DNS records for google.com using nslookup — identify A, AAAA, MX, TXT records

Test connectivity to a website on port 443 using telnet or nc

Use traceroute/tracert to see network path to a remote server

MEDIUM Level

Given VNet 10.0.0.0/16, divide it into 4 equal subnets. Calculate range and usable IPs for each.

Design CIDR plan for an application with: 500 web servers, 100 app servers, 10 database servers. What subnet sizes needed?

Create a DNS zone file with A, CNAME, MX records for a fictional company

Document OSI layers involved when you browse <https://website.com> — what happens at each layer?

Set up a simple HTTP server and test different HTTP methods (GET, POST, PUT, DELETE) — observe status codes

DIFFICULT Level

Design complete IP addressing scheme for multi-region deployment: US (East, West), Europe (North), Asia (Southeast). No overlap, room for growth.

Implement split-horizon DNS: internal clients get private IPs, external clients get

public IPs for same hostname

Troubleshoot a connectivity issue: Given symptoms and network diagram, identify where problem is (Layer 3? Layer 4? Layer 7?)

Calculate optimal MTU for VPN tunnel considering overhead from IPsec, explain impact of fragmentation

Design DNS architecture with high availability: primary/secondary servers, zone transfers, DNSSEC

14.2 Azure Networking — Practice (Use Free Tier)

EASY Level

Create a VNet with address space 10.0.0.0/16 in Azure Portal

Add three subnets: Frontend (10.0.1.0/24), Backend (10.0.2.0/24), Database (10.0.3.0/24)

Create an NSG with rules allowing HTTP (80) and HTTPS (443) from internet

Deploy a simple VM with no public IP, then add a public IP later

Test NSG rules by trying to SSH/RDP with rule disabled vs enabled

MEDIUM Level

Create hub-and-spoke topology: 1 hub VNet, 2 spoke VNets, configure peering

Deploy Internal Load Balancer with 2 backend VMs, configure health probe

Set up Azure Bastion and access a VM without public IP

Create Private Endpoint for Storage Account, verify DNS resolution to private IP

Configure Application Gateway with path-based routing: /api → backend1, /app → backend2

DIFFICULT Level

Implement complete hub-spoke with Network Virtual Appliance (NVA) for spoke-to-spoke traffic

Configure User-Defined Routes (UDR) to force all internet traffic through firewall/NVA

Set up VPN Gateway with Point-to-Site connection, test connectivity from on-prem

Deploy multi-region VNets with Global VNet Peering, implement Traffic Manager for failover

Configure WAF on Application Gateway with OWASP rules, test SQL injection protection

14.3 Mini Project Ideas

Project 1: Secure 3-Tier Application Architecture

Goal — Deploy WordPress with web, app, and database tiers in Azure

Steps:

- Create VNet (10.0.0.0/16) with 4 subnets: Bastion, Web, App, DB
- Deploy Azure Bastion for secure access
- Create NSG for each tier: Web allows 80/443, App allows app port only from Web, DB allows 3306 only from App
- Deploy 2 web server VMs in Web subnet behind Load Balancer
- Deploy 1 app server VM in App subnet
- Deploy Azure Database for MySQL in DB subnet with Private Endpoint
- Configure Application Gateway with WAF in front of web tier
- Test: access WordPress via App Gateway public IP, verify no direct access to backend VMs
- Document network diagram showing traffic flow

Project 2: Multi-Region DR Setup

Goal — Deploy application in two regions with failover

Steps:

- Create VNet in East US and West Europe
- Configure Global VNet Peering between them
- Deploy identical app stack in both regions
- Set up Azure Database replication (geo-replication) primary → secondary
- Configure Azure Traffic Manager with priority routing (East US primary)
- Test failover: shut down East US resources, verify Traffic Manager routes to West Europe
- Measure failover time and document RTO/RPO

Project 3: Secure Hybrid Cloud Connectivity

Goal — Connect on-premises network to Azure securely

Steps:

- Create VNet with GatewaySubnet (/27 minimum)
- Deploy VPN Gateway in VNet
- Simulate on-premises: create separate VNet to act as 'on-prem' with VPN Gateway
- Configure VNet-to-VNet VPN connection
- Set up routing: resources in each VNet can reach other via private IPs
- Test connectivity: VM in VNet A pings VM in VNet B using private IP
- Configure NSG to allow specific traffic between 'on-prem' and cloud
- Document: connection diagram, IP addressing, routing tables

Project 4: Network Monitoring & Troubleshooting Lab

Goal — Set up monitoring and practice troubleshooting

Steps:

- Deploy VNet with 3 VMs: web, app, database
- Enable NSG Flow Logs on all NSGs
- Configure Network Watcher: Connection Monitor, packet capture
- Intentionally create issues: block port in NSG, break routing, misconfigure DNS
- Use Network Watcher to diagnose each issue:
 - IP Flow Verify: test if traffic allowed/denied
 - Next Hop: verify routing paths
 - Connection Troubleshoot: check end-to-end connectivity
- Document each issue and how you diagnosed it

PART 5

COMMON MISTAKES & TROUBLESHOOTING

15. Common Networking Mistakes

15.1 IP Addressing & Subnetting Mistakes

Mistake 1: Overlapping VNet Address Spaces

Problem: Trying to peer VNet A (10.0.0.0/16) with VNet B (10.0.0.0/16).

Error: Peering fails because address spaces overlap.

Fix: Use non-overlapping ranges. VNet A: 10.0.0.0/16, VNet B: 10.1.0.0/16, VNet C: 10.2.0.0/16.

 **WARNING:** Plan IP addressing from the start! Changing VNet address space after deployment is extremely difficult.

Mistake 2: Subnet Too Small

Problem: Created /28 subnet (16 IPs) for AKS cluster or Application Gateway.

Error: Deployment fails — not enough IPs. Azure reserves 5, AKS needs minimum 30+ for scaling.

Fix: Use /24 (251 usable) as default subnet size. For AKS production: /22 or /20.

 **TIP:** Always plan for growth. Better to have unused IPs than run out later.

Mistake 3: Forgetting Azure's Reserved IPs

Problem: Created /27 subnet (32 IPs), deployed 27 VMs, then deployment fails.

Cause: Azure reserves first 4 and last 1 IP in every subnet. /27 has $32 - 5 = 27$ usable.

Fix: Account for 5 reserved IPs when calculating capacity.

15.2 NSG (Firewall) Mistakes

Mistake 1: Wrong Rule Priority

Problem: Created DenyAll rule with priority 100, then AllowHTTPS at 200. HTTPS blocked.

Cause: Lower priority number = higher priority. Rule 100 evaluated first and denies.

Fix: Allow rules should have lower numbers (100-1000), Deny rules higher (4000+).

Mistake 2: Allowing 0.0.0.0/0 on SSH/RDP

Problem: NSG allows port 22 from 0.0.0.0/0 (entire internet).

Security Risk: VM will be scanned and attacked within minutes. Brute-force attempts constantly.

Fix: Use Azure Bastion. OR allow SSH only from specific IPs (corporate VPN range). NEVER allow 0.0.0.0/0 on 22, 3389, 3306.

Mistake 3: Forgetting Inbound vs Outbound

Problem: Created inbound rule allowing 443, but app can't make HTTPS calls to internet.

Cause: Outbound rules are separate. Default allows outbound, but custom NSG might have denied it.

Fix: Check both inbound AND outbound rules. Applications need outbound for API calls, updates, etc.

15.3 Load Balancer / Application Gateway Mistakes

Mistake 1: Using Layer 4 LB for HTTP Routing

Problem: Trying to route /api and /app to different backends using Load Balancer.

Cause: Load Balancer is Layer 4 — doesn't understand HTTP URLs.

Fix: Use Application Gateway for HTTP routing, Load Balancer only for simple port-based.

Mistake 2: Health Probe Misconfiguration

Problem: Application Gateway marks all backends as unhealthy, no traffic flows.

Cause: Health probe endpoint wrong (/health vs /healthz), or app doesn't respond with 200.

Fix: Verify health probe path exists and returns 200 OK. Test with curl from VM.

Mistake 3: Backend Pool Empty

Problem: Created Application Gateway but no traffic reaches backends.

Cause: Forgot to add VMs/VMSS to backend pool.

Fix: Always verify backend pool has targets and they're healthy in monitoring tab.

15.4 DNS Mistakes

Mistake 1: Forgetting to Update DNS After IP Change

Problem: Changed VM IP or deployed new load balancer, but users still get old IP.

Cause: DNS cached at client, recursive resolver, and browser.

Fix: Wait for TTL to expire OR proactively lower TTL to 300 (5 min) before change, make change, then raise TTL back.

Mistake 2: Incorrect DNS for Private Endpoint

Problem: Created Private Endpoint but app still connects to public IP.

Cause: Private DNS zone not linked to VNet, or DNS record missing.

Fix: Verify: nslookup myaccount.blob.core.windows.net resolves to private IP (10.x), not public. Link Private DNS zone to VNet.

Mistake 3: Split-Brain DNS Issues

Problem: Internal users can't resolve external domain, or external users get internal IPs.

Cause: Misconfigured split-horizon DNS — same domain in public and private zones.

Fix: Use different subdomains: internal.company.com (private), www.company.com (public).

15.5 Troubleshooting Flowchart

Network Connectivity Issue — Systematic Approach

Can you ping the IP? → NO? → Layer 3 routing problem

Check routing tables (UDR)

Check VNet peering status

Check VPN/ExpressRoute connection

Can you ping the IP? → YES → Is the port open?

Test with: telnet server-ip 443 or nc -zv server-ip 443

NO? → NSG blocking or service not listening

YES? → Layer 7 application issue

Check NSG Rules

Is port allowed in NSG?

Check both subnet NSG and NIC NSG

Verify priority order — deny rule before allow?

Check Service

Is service running? (systemctl status or equivalent)

Is service listening on correct port? (netstat -tuln)

Check service logs for errors

Check Load Balancer / Application Gateway

Are backends healthy?

Is health probe configured correctly?

Check backend pool has members

DNS Issues?

Does nslookup resolve to correct IP?

Is Private DNS zone linked?

Check /etc/hosts for overrides

PART 6

QUICK REVISION GUIDE

16. Quick Reference Cheat Sheets

16.1 OSI Model Quick Reference

| Layer | Protocols | Focus |
|------------------|----------------|--------------------------|
| 7 - Application | HTTP, DNS, SSH | What user interacts with |
| 6 - Presentation | SSL/TLS | Encryption, formatting |
| 5 - Session | Session mgmt | Connection establishment |
| 4 - Transport | TCP, UDP | Ports, reliability |
| 3 - Network | IP, routing | IP addressing, routing |
| 2 - Data Link | Ethernet, MAC | Local network, switches |
| 1 - Physical | Cables | Physical connection |

16.2 Port Numbers — Memorize These

| Port | Service | Port | Service |
|------|---------|-------|------------|
| 22 | SSH | 443 | HTTPS |
| 80 | HTTP | 3306 | MySQL |
| 53 | DNS | 5432 | PostgreSQL |
| 3389 | RDP | 6379 | Redis |
| 25 | SMTP | 27017 | MongoDB |

16.3 HTTP Status Codes Quick Ref

| Code | Meaning |
|---------------------------|-------------------------|
| 200 OK | Success |
| 201 Created | Resource created |
| 301 Moved Permanently | Permanent redirect |
| 400 Bad Request | Invalid request |
| 401 Unauthorized | Not authenticated |
| 403 Forbidden | Not authorized |
| 404 Not Found | Resource doesn't exist |
| 500 Internal Server Error | Server crashed |
| 502 Bad Gateway | Upstream server down |
| 503 Service Unavailable | Server overloaded |
| 504 Gateway Timeout | Upstream server timeout |

16.4 CIDR Quick Reference

| CIDR | Total | CIDR | Total |
|------|--------|------|----------------------|
| /32 | 1 IP | /24 | 256 IPs (254 usable) |
| /31 | 2 IPs | /23 | 512 IPs |
| /30 | 4 IPs | /22 | 1,024 IPs |
| /29 | 8 IPs | /21 | 2,048 IPs |
| /28 | 16 IPs | /20 | 4,096 IPs |
| /27 | 32 IPs | /16 | 65,536 IPs |
| /26 | 64 IPs | /8 | 16,777,216 IPs |

16.5 Azure Networking Quick Decisions

VNet Sizing

VNet: Use /16 (65k IPs) as standard

Subnet: Use /24 (254 IPs) as default

AKS: Use /22 or /20 for production

Bastion: Use /27 minimum

Load Balancing Decision Tree

HTTP traffic? → Use Application Gateway

Non-HTTP traffic (DB, RDP, SSH)? → Use Load Balancer

Need URL routing or WAF? → Use Application Gateway

Simple TCP load balancing? → Use Load Balancer

Secure VM Access

Best: Azure Bastion (no public IPs, browser-based)

Alternative: VPN Gateway (requires client setup)

Never: Public IP on every VM with NSG (security risk)

16.6 Pre-Interview Checklist

Review these the night before interview:

Explain OSI model layers 3, 4, 7 and their relevance to DevOps

Describe TCP vs UDP — when to use each, give examples

Calculate usable IPs from CIDR: /24, /27, /16
Explain difference between 401, 403, 502, 503, 504
Memorize ports: 22, 80, 443, 3306, 5432, 6379, 3389
Describe Azure VNet, subnets, NSG, how they relate
Explain Load Balancer vs Application Gateway
Describe VNet Peering — is it transitive?
What is Private Endpoint and when to use it?
How does Azure Bastion work? Why use it?

END OF NETWORKING STUDY GUIDE

*Master the fundamentals. Understand Azure services. Practice hands-on.
Success will follow!*